

Bottleneck Detection in Parallel File Systems with Trace-Based Performance Monitoring

Julian M. Kunkel Thomas Ludwig

Institute for Computer Science
Parallel and Distributed Systems
Ruprecht-Karls-Universität Heidelberg

Euro-Par 2008

Outline

- 1 Introduction
- 2 System overview
- 3 Performance Metrics and Statistics
- 4 Results
- 5 Summary

Outline

- 1 Introduction
- 2 System overview
- 3 Performance Metrics and Statistics
- 4 Results
- 5 Summary

Motivation

- A parallel file system should utilize all resources
 - Existing distributed/parallel file systems distribute data/metadata
 - Load imbalance leads to degraded performance
- Several factors lead to variation in I/O performance
 - Hardware capability
 - Access pattern of clients
 - Degraded RAID-Arrays
 - Efficiency of optimizations could vary
 - Throughput of a component could depend on the order of requests

Questions Regarding Load Imbalance

- How could the system or users detect load imbalance?
- Which hardware/software cause the load imbalance?
- Is the application's access pattern the reason of the load imbalance?
- How will the user figure out the application behavior leading to imbalance?

Long-term aim

Maybe the user can modify the code to increase efficiency of the system:

- Access pattern
- Data layout on the servers
- Give hints to the file system

Maybe the file system can rebalance access

Precondition

It is important to monitor the systems behavior

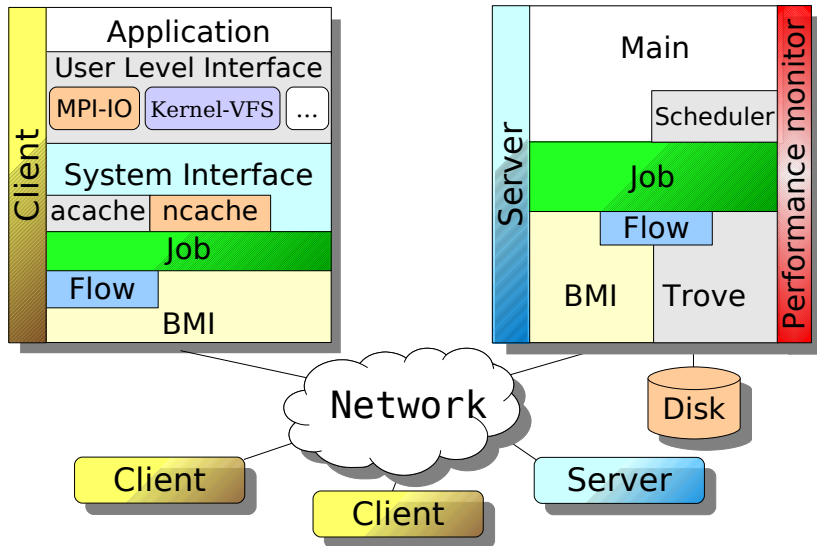
Solution

- Integrate meaningful metrics into the parallel file system
- Allow to query these metrics online
- Visualize the metrics for the users and relate the behavior with the application


Outline

- 1 Introduction
- 2 System overview
- 3 Performance Metrics and Statistics
- 4 Results
- 5 Summary

Architecture of PVFS2



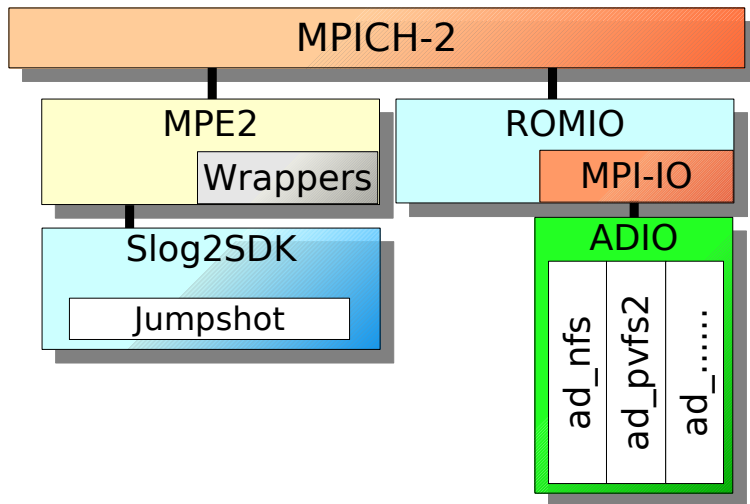
Architecture of the PVFS-Server



Description of the layers/components

- Main
 - Accepts new requests
 - Contains/processes statemachines
 - Scheduler allows concurrent access of non-interfering requests
- Performance monitor
 - Orthogonal layer
 - Stores statistics of internal layers
 - Can be queried remotely
 - Updated in fixed intervals
 - Keeps a history of the statistics

Used Software Environment



PIOviz

- Developed at the university of Heidelberg
- Uses MPE to generate client and server traces
- Add unique ID to PVFS requests in ROMIO
- Provides tools to work on these trace files:
 - Merge
 - Adjust time
 - Correlate client and server activities
- Modifications on Jumpshot:
 - Provide more details on events
 - Allow heights of states proportional to value

Outline

- 1 Introduction
- 2 System overview
- 3 Performance Metrics and Statistics**
- 4 Results
- 5 Summary

Performance Metrics

A metric is a standard unit of measure together with the way it is measured, in order to assess a process, state, or event ...

- Meaning of a component's metric can be different:
 - Represent current state
 - i.e. value is updated instantly - valid only the moment it is fetched
 - e.g. number of currently pending requests
 - Accumulates value during runtime
 - e.g. number of totally processed requests
 - allows to compute values for an arbitrary interval
 - Updated/computed in fixed intervals (statistic)
 - e.g. average number of processed requests

“Absolute” metrics

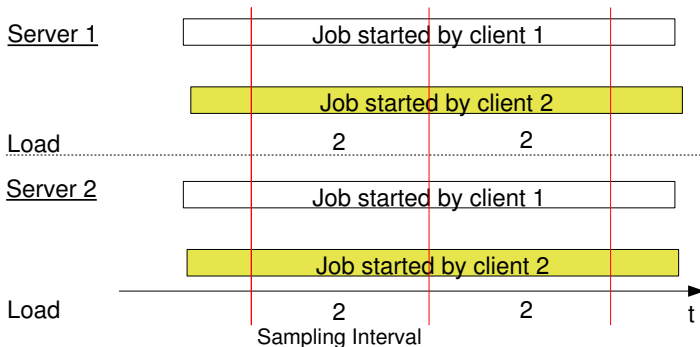
- Measure observable usage or performance
 - e.g. Throughput of network or disk,
 - Processed requests (of a given type),
 - Number of bytes read/written
 - ...
- Is a value of *50 MiB/s* a high value for network throughput?
 - One could relate the value with the maximum network throughput
 - Good value of a server if clients access data only half the time
 - Bad value if the server does not manage to process all requests
 - Is there a congestion in network?
 - Are servers not utilized by client requests?

Relative metrics

- Relate usage/performance with actual demand
 - Idle time of a component in percent within an interval
 - $\gg 0\%$ - Wasted processing capability
 - 0% - Does the component benefit from concurrent operations?
 - Average number of pending jobs within an interval (load-index)
 - The more complex a job the longer it is processed
 - The faster a component operates the shorter the queue
 - e.g. Linux Kernel 60 second system-load
 - e.g. drive queue depth
 - Does the component share its resources among pending operations?
 - Is only a subset of operations serviced at a given time?
 - What if a set of long running jobs is serviced prior short jobs?
 - ...


Balanced Hardware - Long-Running Jobs

- Assume requests which utilize both servers equally




Balanced Hardware - Short-Running Jobs

- Assume a component shares resources among pending requests
- Requests might arrive at different time



Inhomogeneous Hardware - Long-Running Jobs

- Assume server one is twice as fast as server two




Inhomogeneous Hardware - Short-Running Jobs

Server 1

Load

Server 2

Load



Conclusion

Problem

It is not easy to define meaningful metrics!

Expected behavior of a load-index

- Load should be accurate for long-running jobs
- Load average over longer periods should be accurate (even if there are short jobs)

Added Statistics/Metrics (1)

PVFS Statistics

- Request average-load-index
- Flow average-load-index
- I/O subsystem average-load-index
- I/O subsystem idle-time [percent]
- Network average-load-index

Added Statistics/Metrics (2)

Kernel Statistics

- Average kernel load for one minute
- Memory used for I/O caches [Bytes]
- CPU usage [percent]
- Data received from the network [Bytes]
- Data send to the network [Bytes]
- Data read from the I/O subsystem [Bytes]
- Data written by the I/O subsystem [Bytes]

Relation of several statistics could reveal the component causing imbalance

Outline

- 1 Introduction
- 2 System overview
- 3 Performance Metrics and Statistics
- 4 Results**
- 5 Summary


Test Environment

- 10 node cluster each node equipped with:
 - Two Intel Xeon 2GHz CPUs
 - 1000 MiB memory
 - GBit Ethernet (throughput: 117 MiB/sec)
 - IBM Hard disk (sequential throughput: ≈ 45 MiB/sec)
 - RAID Controller with two disks (≈ 90 MiB/sec)
- Test program
 - Allows to select a level of access ((non-)contig., (non-)collective)
 - Clients write a fixed amount of data
 - Barrier
 - Clients read (their) data
- Disjoint clients and server partition
- Five clients, one metadata server and three data servers


Experiments

- Independent contiguous I/O (each client accesses 200×10 MiB)
- Collective, non-contiguous I/O (each client accesses 4×500 MiB)
 - Note: Size of ROMIO's collective I/O buffer is 4 MiB
- Both cases measured also on an inhomogeneous I/O subsystem
 - One server uses its system disk and not the RAID system

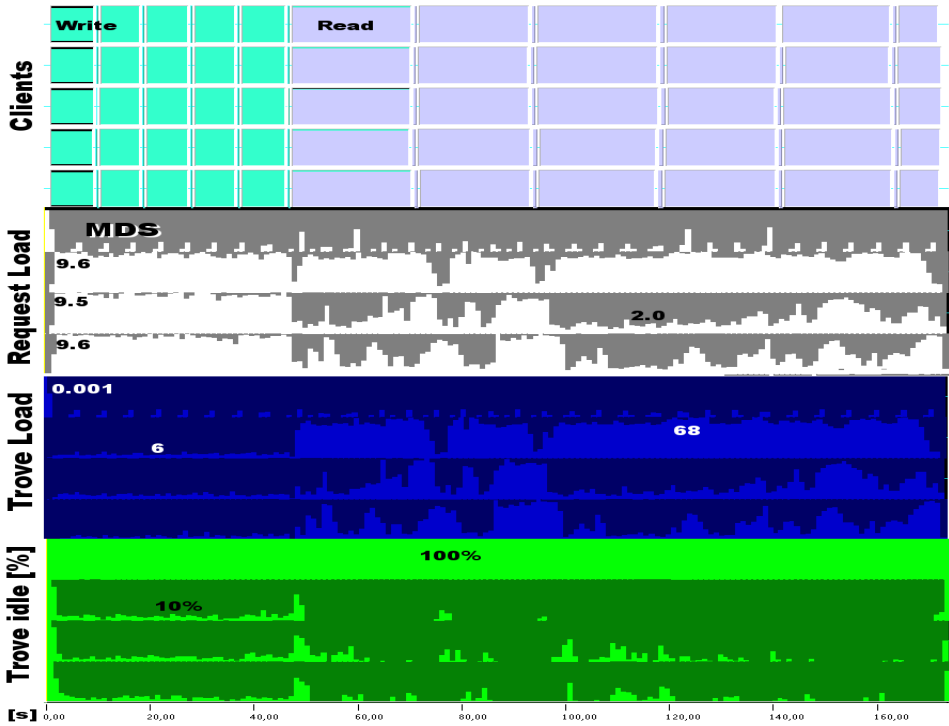
Collective, non-contiguous I/O with homogeneous hardware


Clients**Request Load****Trove idle
[%]****Bytes written
[MB/s]****Bytes read
[MB/s]****t [s]**


Collective, non-contiguous I/O with inhomogeneous hardware




Independent, contiguous I/O







Independent, contiguous I/O with inhomogeneous hardware



	Request load			BMI load			Trove load			Trove idleness [%]		
level 0	7.8	5.1	5.5	2.1	2.1	2.1	36.7	17.9	20.6	5.4	9.7	9.8
level 1	5.1	5	5.1	1.7	1.7	1.7	9.2	9	9.2	26.9	27.6	27.3
level 2	8.7	8.8	8.2	2.2	2.3	2.2	54.4	54.9	49.9	3.4	2.1	4.4
level 3	2.9	3	2.9	1.2	1.2	1.2	4.4	4.8	4.6	56.7	54	56.3
level 0 - 1 datafile	3.2	3.1	3.9	1.2	1.2	1.2	14.8	14.5	20.2	15.5	16.9	9.5
level 1 - 1 datafile	2.7	2.5	2.6	1.1	1.1	1.1	8.5	7.6	7.9	30.3	34.9	33.2
level 2 - 1 datafile	4	4.3	4.1	1.2	1.4	1.2	25.7	28.6	26.6	10.1	5.3	8.8
level 3 - 1 datafile	1.5	1.5	1.5	0.9	0.9	0.9	4.1	4.1	4.1	60.8	60.4	60.9
level 0-inh. I/O	8.4	2.3	2.2	1.2	1.3	1.3	44.1	5	4.6	2.6	40.6	42.1
level 1-inh. I/O	5.8	3.5	3.6	1.2	1.3	1.3	11.9	6	6.2	10.1	49.5	49.2
level 2-inh. I/O	9.2	4.1	4.1	2.6	1.1	1.1	65.2	25.6	25	1.6	39.1	39.8
level 3-inh. I/O	3.5	2.3	2.3	0.9	1	0.9	7.2	3.5	3.5	40.4	63.7	64.3

Outline


- 1 Introduction
- 2 System overview
- 3 Performance Metrics and Statistics
- 4 Results
- 5 Summary**

Summary

- It is important to monitor components performance
 - Relative metrics allow to assess usage with delivered performance
- Introduced an environment which allows to relate server statistics with (MPI) client activities
- Trace could assist the user to detect inefficient MPI-I/O calls and potential reasons
- Relation of several metrics allows to localize the component causing the load imbalance
- More work is needed to assess observed behavior

Optional slides

Unbalanced client access pattern



10 Clients and 10 Serves in PC-Pool

